

Event-Triggered Saturating Attitude Controller

Problem Statement

Diogo Almeida

February 12, 2014

The following is a draft of the problems and proposed solutions for my thesis project, namely the control of a quadcopter attitude using event triggered strategies. I propose the implementation of a saturating controller that fully exploits the available control torques, together with a Lyapunov based event triggering function, [1] and [2] respectively.

1 The attitude controller

Proposed by Fritsch, Henze and Lohmann in [1], the attitude controller implements a control law that stabilizes a quadcopter attitude in two steps. On a first part, the thrust axis of the quad is aligned with the reference direction and, on a second step, the yaw angle of the vehicle is corrected. The control variables here are the displacement angle of the thrust axis, φ , and the yaw error, ϑ . This has advantages if the primary concern of the higher level control revolves around the translational movement of the robot. Another advantage of this controller is its ability to take advantage of the actuating power available by explicitly modelling the maximum torques available to the system and exploiting that knowledge by saturating the control torques whenever possible. This is achieved by adopting an energy shaping approach, where a desired energy for the system is designed and a damping strategy that penalizes movements that go against the reference equilibrium, while boosting the ones that go in the right way, is applied. The resulting control torques have the expression

$$\boldsymbol{\tau}(\mathbf{x}) = \mathbf{T}(\mathbf{q}) - \mathbf{D}(\mathbf{x})\boldsymbol{\omega} \quad (1)$$

Where \mathbf{q} is the quaternion that represents the attitude error of the quadcopter, $\boldsymbol{\omega} = [\omega_x \ \omega_y \ \omega_z]^\top$ are the angular velocities of the quad around its three body frame axis, $\mathbf{T}(\mathbf{q})$ is the torque field that the desired energy would generate and $\mathbf{D}(\mathbf{x})$ is a positive semi-definite damping matrix that allows for the control saturation as well as the almost global asymptotic stability of the controlled system. Finally, the state of the system with respect to the attitude is given by $\mathbf{x} = [\mathbf{q} \ \boldsymbol{\omega}]^\top$.

The control law (1) is designed to ensure that a Lyapunov function (LF) consisting in the sum of the system kinetic energy with an artificial potential




Figure 1: The applied torques and resulting behavior of the control variables for the saturating controller [1]

energy has a time derivative given by

$$-\boldsymbol{\omega}^\top \mathbf{D}(\mathbf{x}) \boldsymbol{\omega} \leq 0 \quad (2)$$

and the authors of [1] were able to prove asymptotic stability based on LaSalle's Stability theorem. The baseline result for this controller, under the initial conditions in [1] are given in figure 1.

2 Event-triggered strategy

2.1 Lyapunov based scheduling

The main problem in event-triggered implementations is to find an execution rule that allows to generate the sampling instants t_k in such a way that the stabilization of the system is ensured, while avoiding accumulation points. One systematic approach to the problem is given by [2], where a sampling instant is generated every time the time derivative of the LF, $V(\mathbf{x})$, of the system violates an inequality condition. This strategy requires the LF to be bounded above and below by strictly increasing functions of the state and for its time derivative to be upper bounded by an 'error term' $-\alpha(|\mathbf{x}|) + \gamma(|\mathbf{e}|)$, where \mathbf{e} is the difference between the state at the current time t and the last sampling instant t_k

$$\begin{aligned} \underline{\alpha}(|\mathbf{x}|) &\leq V(\mathbf{x}) \leq \bar{\alpha}(|\mathbf{x}|) \\ \dot{V}(\mathbf{x}) &\leq -\alpha(|\mathbf{x}|) + \gamma(|\mathbf{e}|) \end{aligned} \quad (3)$$

and, by ensuring

$$\gamma(|\mathbf{e}|) \leq \sigma \alpha(|\mathbf{x}|), \quad 0 < \sigma < 1 \quad (4)$$

we get $\dot{V}(\mathbf{x}) \leq 0$ at all times.



2.2 Threshold based scheduling



Following [3], a simpler alternative to use as execution rule is to define a threshold $\bar{\epsilon}$ for the difference between the control variable at the time t_k , *i.e.* the last time the control signal was updated, and the current value at the time t , over which the control signal is recomputed. It doesn't offer the guarantee that the system will remain stable when enforcing this update rule, and as such it probably needs to be tuned experimentally to obtain acceptable results without degrading performance too much.

3 Possible approaches



To generate an event rule, we have several alternatives, from designing a Lyapunov based rule to refreshing the control when the error with respect to the reference is bigger than a threshold. There are stability as well as complexity issues that must be analysed.



3.1 Lyapunov based Saturating Event Triggered control

The approach that motivates me the most is the one discussed in [2], and as such I would like to implement the saturating controller, together with an event generator where the event triggering rule is obtained from the time derivative of the LF of the controlled system, (2). For that, I need to find an $\alpha(|\mathbf{x}|)$ and $\gamma(|\mathbf{e}|)$ class- \mathcal{K} functions that comply with (3) and allows for the definition of the update rule (4). Fritsch et al. [1] define $\mathbf{D}(\mathbf{x})$ as

$$\mathbf{D}(\mathbf{x}) = \begin{bmatrix} \kappa_{xy}(\mathbf{x})\mathbf{D}_{xy}(\mathbf{x}) & \mathbf{0} \\ \mathbf{0} & \kappa_z(\mathbf{x})d_z(\mathbf{x}) \end{bmatrix} \quad (5)$$

where

$$\mathbf{D}_{xy}(\mathbf{x}) = \frac{d_\varphi(\mathbf{x})}{1 - q_p^2} \begin{bmatrix} q_x^2 & q_x q_y \\ q_x q_y & q_y^2 \end{bmatrix} + \frac{d_\perp}{1 - q_p^2} \begin{bmatrix} q_y^2 & -q_x q_y \\ -q_x q_y & q_x^2 \end{bmatrix}$$

hence, the time derivative of the LF, (2) becomes

$$-\frac{\kappa_{xy}(\mathbf{x})}{1 - q_p^2} [\omega_x^2 (d_\varphi(\mathbf{x})q_x^2 + d_\perp q_y^2) + \omega_y^2 (d_\varphi(\mathbf{x})q_y^2 + d_\perp q_x^2)] - \omega_z^2 \kappa_z(\mathbf{x})d_z(\mathbf{x}) \quad (6)$$



one can easily see that, since in the equilibrium $q_p = \pm 1$, by multiplying the first element in the sum (6) by $1 - q_p^2$, we obtain a function that is always bigger than (6), that can be used to create the event rule (4) while ensuring $\dot{V}(\mathbf{x}) \leq 0$ for all \mathbf{x} different than $\mathbf{0}$.



3.2 Heuristic approach

Another alternative is to discard altogether the Lyapunov based approach and try to generate an event-triggering mechanism based on either the state error

as defined above, or the error with respect to the reference. The first case updates the control signal every time the state deviates more than a certain value from the state at the previous sampling time [3]. In the second case, one can do periodic sampling when outside a 'tolerance region' specified by the system designer [4]. More, if the purpose of the event triggering system is to save network resources, as in [3], the control algorithm may run periodically over an input error that is updated according to the update rule, or the control algorithm may be limited to run only at event times, and the control signal is kept constant in the meanwhile. All these approaches lack stability results *a priori* and require an analysis case by case.

4 Lyapunov based event triggering results

This is the approach proposed in [2]. Given a controller that stabilizes the system, if we fulfil the conditions in the theorems presented, we can get a control update rule that preserves the stability of the system. The main problem with this approach is that it is not easy to find an $\alpha(|\mathbf{x}|)$ that is strictly increasing and satisfies the inequality. Remembering that a desired equilibrium has two corresponding points in the quaternion space, $\boldsymbol{\omega} = \mathbf{0}$ and $\mathbf{q} = [0 \ 0 \ 0 \ \pm 1]$, we have that

$$\alpha(\mathbf{x}) = -\boldsymbol{\omega}^\top D_\alpha(\mathbf{x}) \boldsymbol{\omega} \quad (7)$$

where

$$D_\alpha(\mathbf{x}) = \begin{bmatrix} \kappa_{xy}(\mathbf{x})(1 - q_p^2) \mathbf{D}_{xy}(\mathbf{x}) & \mathbf{0} \\ \mathbf{0} & \kappa_z(\mathbf{x}) d_z(\mathbf{x}) \end{bmatrix}$$

satisfies (3), though it is not a class- \mathcal{K} function, nor a function of the norm of the state. From what I understand, the need for the function to be strictly increasing with the norm of the state is to ensure that there are no accumulation points. This can be solved by imposing a minimum inter sampling time¹. The stability of the resulting system is ensured by (4). Indeed, the implementation of the update rule with $\alpha(\mathbf{x})$ as defined in (7) allows for the convergence of the controlled variables without any noticeable performance issues, see figure 2 for the results with $\sigma = 0.9$.

4.1 How to deal with accumulation points

The main problem with this result, in my view, is the decrease in the inter-sampling time as the simulation goes on. Eventually, when the state of the system stabilizes, inter-sampling time coincides with the sampling time I use for the simulation. It seems as if $\gamma(|\mathbf{e}|)$ is decreasing slower than $\alpha(\mathbf{x})$, and this results in the inequality (4) being violated every time it is checked. I assume this is a direct consequence of not enforcing $\alpha(\mathbf{x}) \in \mathcal{K}$. My proposal is to add a very small ϵ_x to (4) so as to become

$$\gamma(|\mathbf{e}|) \leq \sigma \alpha(|\mathbf{x}|) + \epsilon_x, \quad 0 < \sigma < 1 \quad (8)$$

¹But will that remove stability insurances?



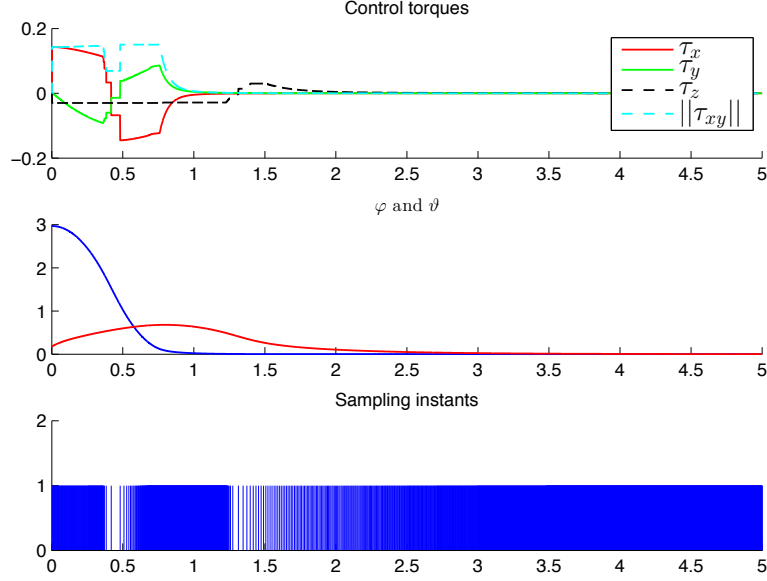


Figure 2: Results using the update rule (4) with $\alpha(\mathbf{x})$ defined as in (7)



With this new update rule we introduce a dead zone for the state-error that may induce stability problems, but at the same time allows for the inter-sampling time to grow unbounded if the state is in the desired equilibrium, see figure 3 with $\sigma = 0.9$ and $\epsilon_x = 0.05$.

5 Error based methods

Besides tying the event triggering to the evolution of the time derivative of the LF of the system, one can try to implement rules based purely on the state error, that is, on how much has the state changed since the last time the control was updated or even by defining an area of admissible error, where the control is kept constant while the state of the system remains inside it.



5.1 State error method

Here, the update rule is based on the evolution of the state since the last time the control was updated. Given the error

$$\mathbf{e} = \mathbf{x}(t_k) - \mathbf{x}(t)$$




Figure 3: Results for the execution rule (8)

the control is updated every time the norm of the state error grows beyond a certain threshold

$$|\mathbf{e}| > \sigma_e \quad (9)$$

One advantage of this rule relies on its simplicity. If the state is diverging too much from the one we had the last time the control was updated, then it is time to do it again. Stability may be a concern here, at least when near the equilibrium, the control may be such that the state doesn't stay where it is meant to be, resulting in poor performance and extra control updates. The value of the threshold σ_e impacts the results significantly as well. For $\sigma_e = 0.1$ the observed results are the ones in figure 4. With this threshold value, it is noticeable that the controller lets the state move away from the desired equilibrium and reacts later in the simulation to correct it. This results in a diverging behavior of the state when it should be resting in the reference value (figure 5). A similar behavior is observed when using the rule (8) for larger values of σ_e , cases where there are less updates to the control signal, but results in larger tolerance for positive $\dot{V}(\mathbf{x})$.



5.2 Error with respect to the reference

Another simple idea is to update the control value every time the state leaves an admissible set [4]. This is similar to having a relay in the control loop that

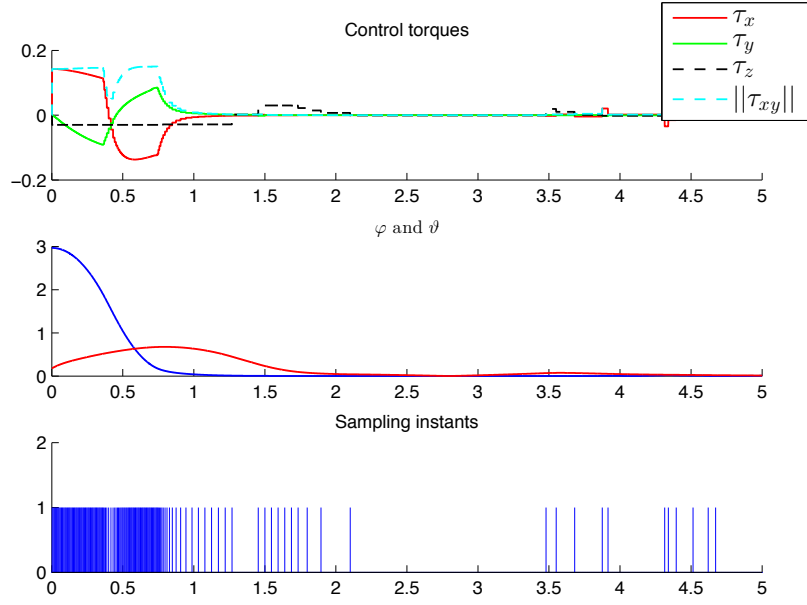


Figure 4: Results for the execution rule (9)




Figure 5: Details of the diverging behavior when using execution rule (9)




Figure 6: Results for the admissible set rule

closes the feedback loop only when the error grows above a certain threshold. Here, the error may be given as

$$\mathbf{e} = \mathbf{x}(t) - \mathbf{x}_d$$

where \mathbf{x}_d is the state in the equilibrium. Since the proposed controller works on the error between the current state and the reference, the desired equilibrium is always the same, which makes it easy to implement the update rule. In figure 6 are depicted the results for an admissible deviation of 1 percent of the state norm. For every value of the admissible error experimented with, the oscillating behavior appeared, sooner or later. That seems reasonable, since the system is unstable while in open-loop. Changing the error value has an impact in the overall performance of the system while near the equilibrium, as well as on the number of control updates while near the admissible set.



6 Comparison between approaches

With the exception of the update rule (4)², all the approaches analysed here do not ensure that the asymptotic stability of the system is preserved. That may

²Or even with the inclusion of the update rule (4). Is the Lyapunov based analysis (3) only valid for continuous systems?



be ok or lead to unacceptable behavior. There may be a tradeoff between performance and number of time events. Finally, some techniques may be objectively better than others in some aspects. Some comparison rules are proposed to benchmark the event rules.

6.1 Number of updates

The trivial measurement of performance is the number of updates of the control signal during the same time period. Less updates means that the control algorithm is called less often, resulting in computational savings. The simulation results for the cases above are presented in table 1.

| Strategy | Updates |
|---------------------|------------|
| Baseline | 5000 |
| Rule (4) | 2801 (56%) |
| Rule (8) | 439 (9%) |
| Rule (9) | 155 (3%) |
| Admissible set rule | 2109 (42%) |

Table 1: Number of control updates

6.2 Control signal energy

One of the purposes of using event-triggered control is to reduce the energy consumption of the overall system, by minimizing the CPU usage. Those gains may be lost if the resulting control signals uses more energy. To quantify the differences in the torques energy, I use

$$E_i = \frac{1}{T} \sum_{t_k=0}^T \tau_i^2, \quad i = \{x, y, z\}$$

the respective results are presented in table 2.

| Strategy | τ_x | τ_y | τ_z |
|---------------------|---------------|----------------|----------------|
| Baseline | 2.3333 | 0.53697 | 0.25928 |
| Rule (4) | 2.4573 (105%) | 0.5833 (109%) | 0.26245 (101%) |
| Rule (8) | 2.477 (106%) | 0.60163 (112%) | 0.26369 (102%) |
| Rule (9) | 2.3949 (102%) | 0.54307 (101%) | 0.29017 (112%) |
| Admissible set rule | 2.539 (108%) | 0.54029 (101%) | 0.28066 (108%) |

Table 2: Torques energy

7 Work plan

There are many things I have not looked into, or have approached only superficially. The parameter I used for each approach were not tuned to obtain the best results. I just used ones that 'seemed to make sense'. This may have a significant impact in the comparisons I did. I have not tried to look deeply into the definition of (5) to try to find a better $\alpha(\mathbf{x})$ function, or even one that fulfills the theorems in [2]. Furthermore, I am only applying these approaches in simulations where I assume full access to the system state. I could simulate the sensors and respective sensor noise. Maybe that will reveal problems with some of the techniques above. Another issue concerns the implementation. Getting to know the system and the low level control code that is already implemented in the Arducopter may take some time. There is a need to estimate the maximum available torques that the system can output, to ensure that the proposed controller will behave properly. Finally, setting up a systematic way of testing and comparing these approaches in the real system is something I would also like to do.



References

- [1] O. Fritsch, B. Henze, and B. Lohmann, "Fast and saturating attitude control for a quadrotor helicopter," 2013.
- [2] P. Tabuada, "Event-triggered real-time scheduling of stabilizing control tasks," 2007.
- [3] D. Lehmann and K. H. Johansson, "Event-triggered pi control subject to actuator saturation," 2012.
- [4] K. J. Aström, "Event based control," in *Analysis and Design of Nonlinear Control Systems* (A. Astolfi and L. Marconi, eds.), pp. 127–147, Springer Berlin Heidelberg, 2008.